

UNIVERSITY *of* WASHINGTON

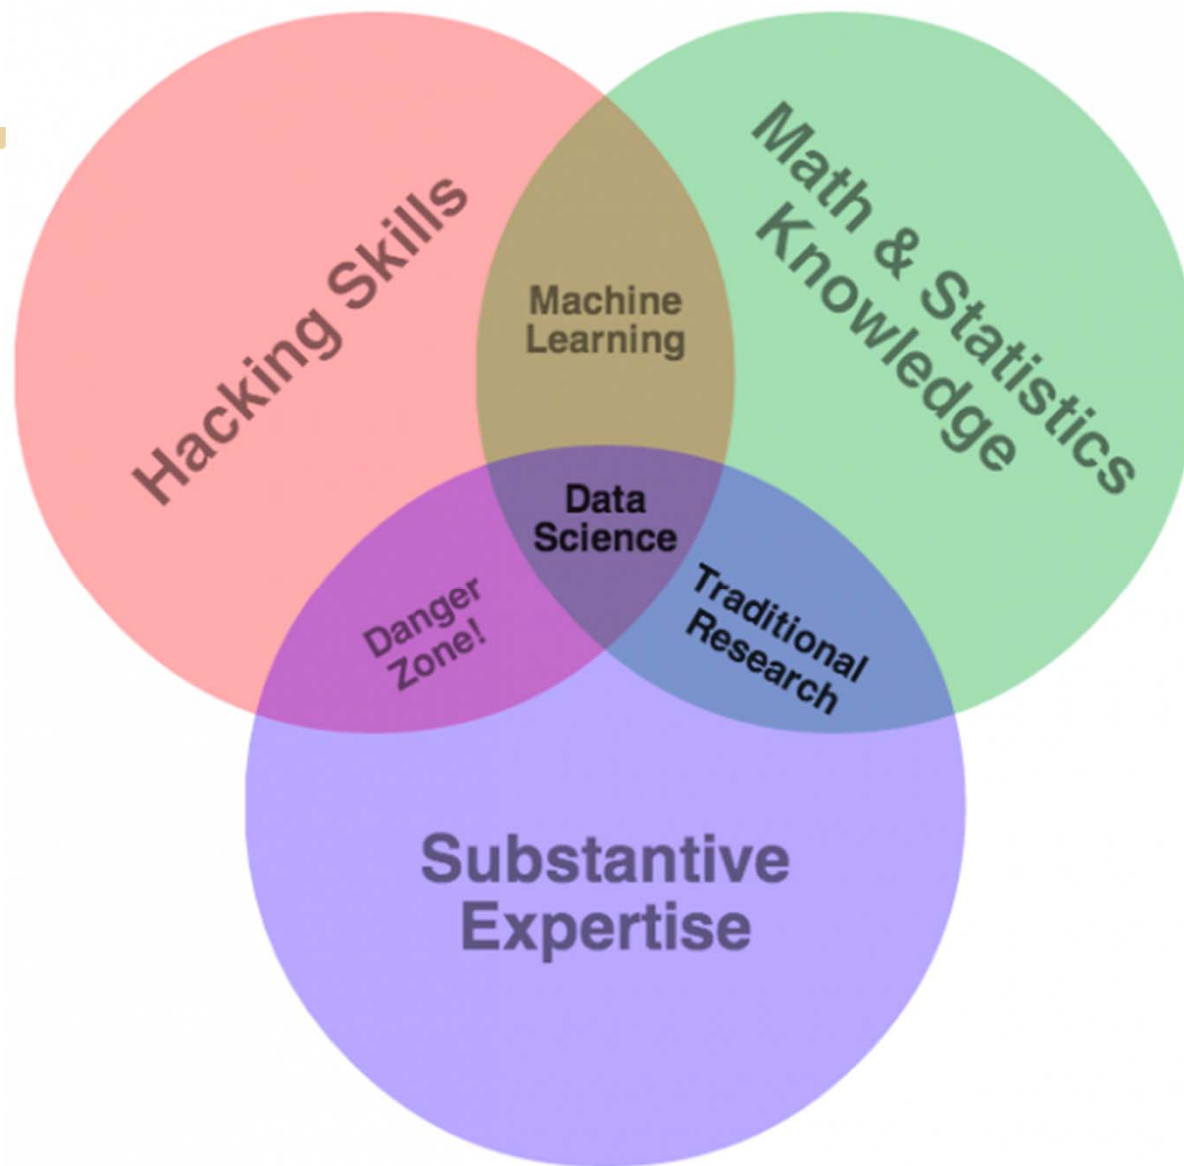
# Data Science UW

## Methods for Data Analysis



Introduction and Data Exploration  
Lecture 1  
Stephen Elston





**W**

# Course Purpose

- > This course focuses on essential concepts
- > We are building foundations for your data science skills
- > Course Objectives:
  - Learn methods to explore and understand data.
  - Understand the core concepts of statistics.
  - Understand and implement various statistical procedures in R.
  - Describe and interpret analytical results from common statistical methods.
  - Expand R programming skills to be able to write/test/log code from scratch.
  - Work with structured and unstructured data.

- > See syllabus for more information:

- <https://canvas.uw.edu/courses/1087732/pages/course-abbrev-course-syllabus>



# Course Requirements and Grading

This course will be graded by attendance, homework, and an individual project.

- > Attendance: You MUST attend at least 8 out of 10 classes. **This is a non-negotiable UW requirement.**
- > Homework must be completed by the start of the next class. (Assigned weeks 1-8).
  - Returned as a 0,1, or 2.
    - > 0 = Not done or a major part wrong/missing.
    - > 1 = Completed, but missing or got wrong 1 or 2 parts.
    - > 2 = Completed with at most minor issues. Demonstrates full understanding of subject.
- > Individual Project: Due at the start of the last class.
  - Counts as 8 points.



# Course Requirements and Grading

There is a total of 24 possible points. (16 pts for hmk + 8 project)

- > Must get 18 total points to pass.
- > All homework assignments must use good R coding technique
- > The individual project must be production level code.



# Office Hours and Contact Information

---

## > Contact me at:

- [stephen.elston@quantia.com](mailto:stephen.elston@quantia.com)

## > When I'm *usually* available:


- Off/on for simple things during work. (M-F 8am-5pm PST)
- Sunday various afternoon/evening times.

Emergency contact: 402-980-3192




# Review

## PYTHON




## JAVA



## C++



## UNIX SHELL



## ASSEMBLY



## C



## LATEX



## HTML



# SQL Review

- > SQL is the 'Lingua Franca' of data access
- > SQL (to know):
  - *Create tables*
  - *Drop tables*
  - Select, where, groupby
  - Joins (Inner, outer, right, left)
  - Temp tables
  - Coalesce, Cast, Case





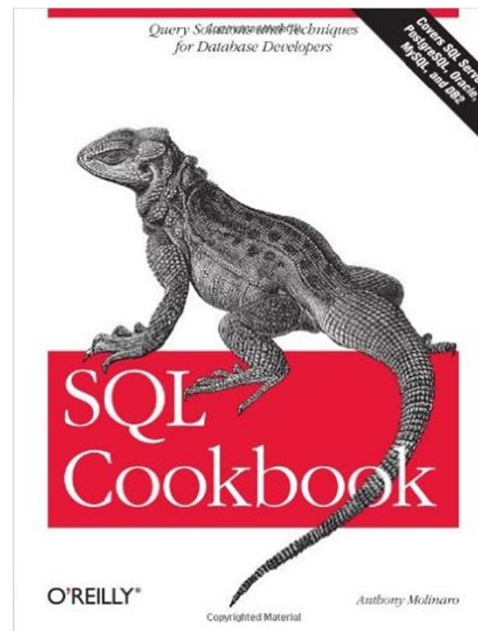
# SQL Resources

## SQL Tutorial and Resources

<http://www.w3schools.com/sql/>

## Querying with Transact SQL Course, Graeme Malcom

<https://www.edx.org/course/querying-transact-sql-microsoft-dat201x-3>



# Prepare for R Demos

---

> Install R

<https://cran.r-project.org/>

-or-

<https://mran.revolutionanalytics.com/download/>

> Install RStudio

<https://www.rstudio.com/products/rstudio/download/>



# GitHub

---

- > Code, data and slides for this course are in a GitHub repository

<https://github.com/StephenElston/DataScience350>

- > Install GitHub for desk top

<https://help.github.com/desktop/guides/getting-started/installing-github-desktop/>

- Or, just download the zip files



# R Review

## > R resources:


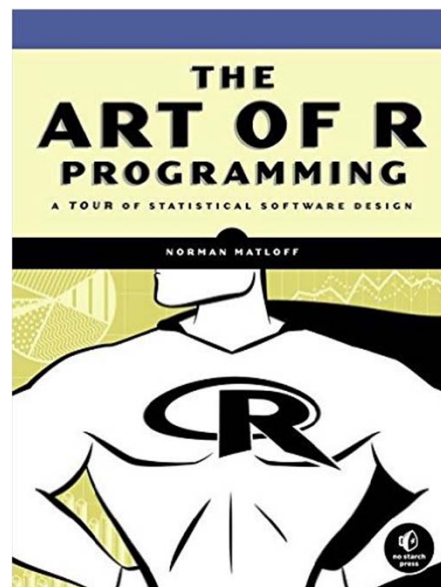
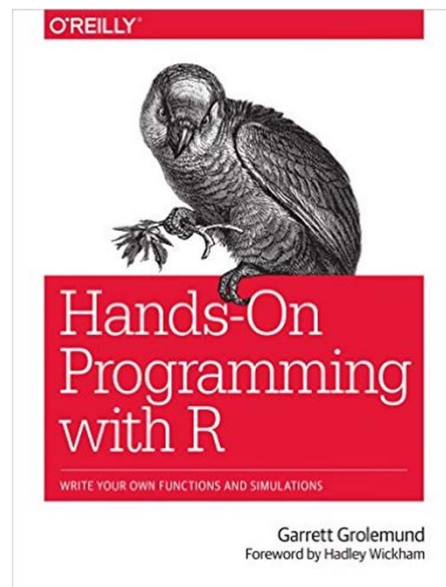
- R page:
  - > <http://www.r-project.org/other-docs.html>
- Stackoverflow:
  - > <http://www.stackoverflow.com>
- ‘Little’ R intro:
  - > <http://cran.r-project.org/doc/contrib/Rossiter-RIntro-ITC.pdf>
- Quick R:
  - > <http://statmethods.net/>
- There are many tutorials available online, e.g.,
  - > <http://cyclismo.org/tutorial/R/>
- Google’s Style Guide:
  - > <http://google-styleguide.googlecode.com/svn/trunk/google-r-style.html>



# More R Resources

## R Inferno, Pat Burns

[http://www.burns-stat.com/pages/Tutor/R\\_inferno.pdf](http://www.burns-stat.com/pages/Tutor/R_inferno.pdf)



W

# Statistics Review

## > Familiar Concepts:

- Discrete vs. Continuous Distributions
- Probability
- Statistics
- $y = mx + b$  vs  $\bar{Y} = \mathbf{M} \cdot \bar{X} + \mathbf{B}$

> These concepts are the focus of this course.



# Counting Review

## > Factorials

- Count # ways to order N things =  $N!$

## > Permutations

- Count # of ways to **order** R things from N things =  $N!/(N-R)!$
- Ordering matters
- $P(N,R)$

## > Combinations

- Count # of ways to **group** R things from N things =  $N!/(R!(N-R)!)$
- Ordering doesn't matter
- $C(N,R)$  or  $\binom{N}{R}$

> We will talk about this in depth next class.



# Data Distributions (Discrete)

## > Discrete Distribution Properties

- Sum of probability of all possible events must equal 1.
- Probability of event equal to value of distribution at point.
- All values strictly in range 0-1.





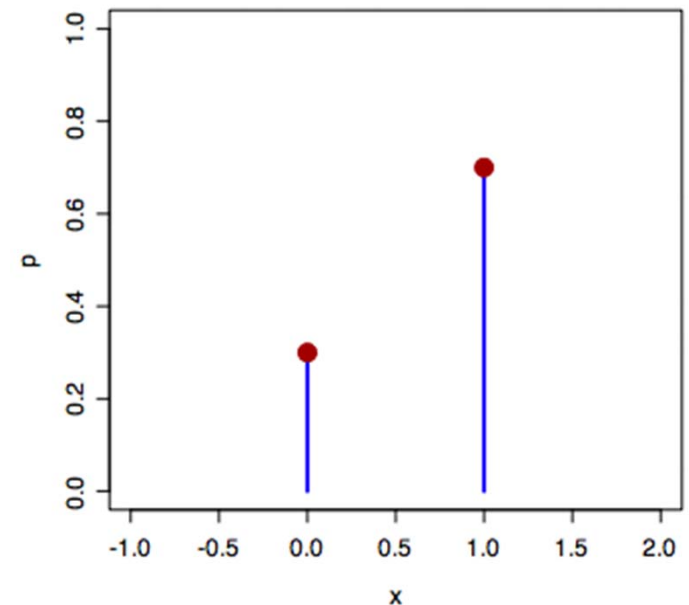
# Data Distributions (Discrete)

> Bernoulli (1 event, e.g.: coin flip)

$$P(x) = \begin{cases} p & \text{if } x = 1 \\ (1 - p) & \text{if } x = 0 \end{cases}$$

$$P(x) = p^x (1 - p)^{(1-x)} \quad x \in \{0,1\}$$

- Mean =  $p$
- Variance =  $p(1-p)$



# Data Distributions (Discrete)

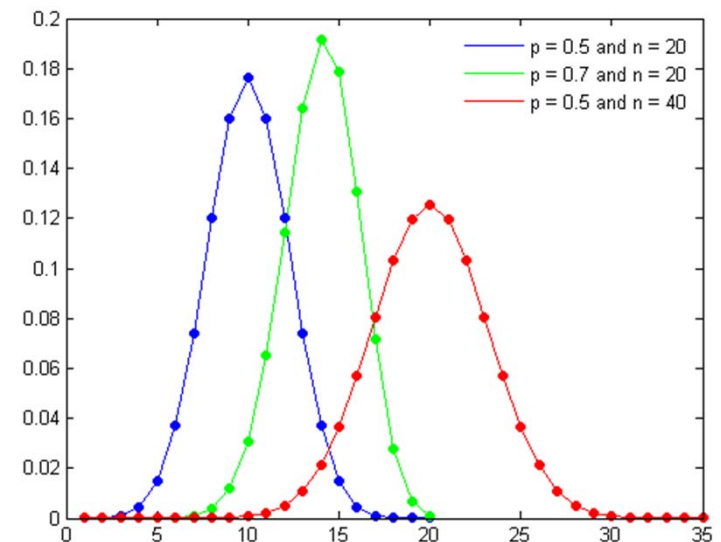
## > Binomial (Multiple Bernoulli's Events)

- Multiple Independent events = Product of Bernoulli Probabilities

$$P(x|N, p) = \binom{N}{x} p^x (1 - p)^{(N-x)}$$

- Mean =  $np$
- Variance =  $np(1-p)$

Note: for larger  $n$ , we approximate this by a normal distribution.




# Data Distributions (Discrete)

> Poisson (Count of number of events in a time span)

$$P(x|\lambda) = \frac{\lambda^x}{x!} e^{-\lambda}$$

- Mean =  $\lambda$
- Variance =  $\lambda$

Interpret as the rate of occurrence of an event is equal to lambda in a finite period of time.





# R Demo



Discrete distributions



# Data Distributions (Continuous)

## > Continuous Distribution Properties

- Area under the curve must be equal to 1.
- Probability a range of values of an event equal to AREA under curve.
- No negative values.
- Probability of a single, exact value is 0.



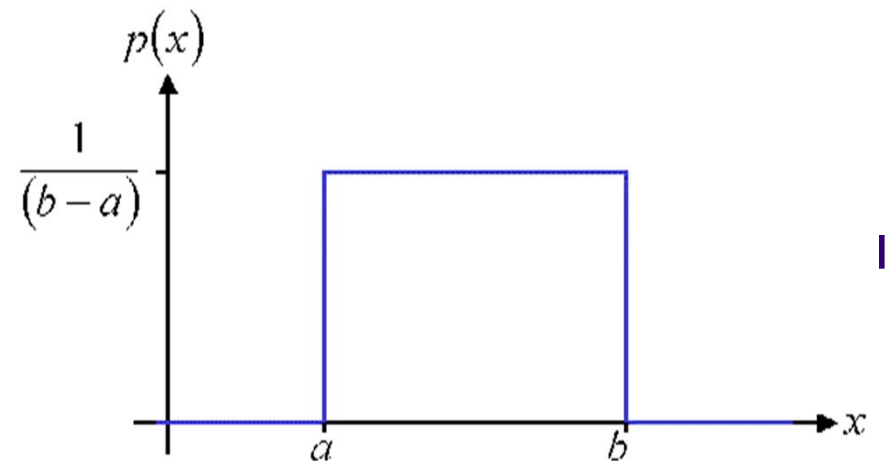
# Data Distributions (Continuous)

> Uniform (flat, bounded)

$$P(x) = \begin{cases} \frac{1}{(b-a)} & \text{if } a \leq x \leq b \\ 0 & \text{if } x < a \text{ or } x > b \end{cases}$$

> Used for parameter priors. (future discussion)

- Mean =  $(a+b)/2$
- Variance =  $(1/12)(b-a)^2$




# Data Distributions (Continuous)

## > Normal (Gaussian) distribution

- Most common and occurs naturally.
- Defined by a mean and variance only. (standard =  $N(0,1)$ )

$$P(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- Has very nice properties.
- Tests for normality are very important.



# Data Distributions (Continuous)

- > **Student's T** (normal for small samples)
  - Important for hypothesis testing smaller sample sizes.
  - Used for:
    - > Testing of mean value when st. dev. is unknown.
    - > Testing difference between two distribution means.
  - Looks very similar to the normal distribution.







# R Demo



Continuous distributions



# R review and summary statistics

- > Purpose: To gain a clear understanding of your data.
  - How large is it?
  - What columns are of interest?
  - Missing data?
  - Outliers?



# R Vectors, Arrays and Lists

---

- > Vectors have one dimension, one data type
- > Vectors are the atomic object in R
- > R is optimized for vector operations
- > Array has multiple dimensions, all data of one type
- > Matrix is a 2D array
- > Lists are comprised of other R object
- > Elements of a list can be of any type and dimension



# Data frames

---

- > Rectangular tables: cannot be ragged
- > List with special attributes
- > Each column is vector of single type
- > Columns can have names
- > Rows can have names
- > Subsetting function []
- > Column function \$



# Functional Programming with R

## Definition from Wikipedia

In computer science, **functional programming** is a programming paradigm—a style of building the structure and elements of computer programs—that treats computation as the evaluation of **mathematical functions** and **avoids changing-state** and mutable data. It is a declarative programming paradigm, which means **programming is done with expressions or declarations**<sup>l</sup> instead of statements. In functional code, the **output value of a function depends only on the arguments** that are input to the function, so calling a function  $f$  twice with the same value for an argument  $x$  will produce the same result  $f(x)$  each time. **Eliminating side effects**, i.e. changes in state that do not depend on the function inputs, can make it much easier to understand and predict the behavior of a program, which is one of the key motivations for the development of functional programming.



# Functional programming in R

- > R is a functional language
- > Expressions and function are objects
- > Functions can be named or anonymous
- > Use functional operators rather than loops



# R programming practice

- > Vectorize for performance
- > Use functions; **avoid repetitive code!**
- > Make use of functional operators
- > Comment your code
- > Names should mean something





# R Demo



Basic data wrangling and functional programming





# Data Exploration (Descriptive Statistics)

- > Purpose: To gain a clear understanding of your data.
  - How large is it?
  - What columns are of interest?
  - Missing data?
  - Outliers?




# Numerical Exploration

- > `str()`: structure of the data frame
- > `summary()`: summary of each of the columns
- > `head()` / `tail()`: top / bottom of data frame
- > `table()`: frequency table




# Numerical Exploration

- > Quartile ==  $\frac{1}{4}$  values
- > IQR(): inner quartile range ( $Q3 - Q1$ )



# Numerical Exploration

- > `quantile()`: quantiles of numerical vectors
  - Quantiles are inverse values of the CDF (cumulative distribution function).
  - Standard Normal: (shown in figure)
    - >  $\text{Quantile}(0.5) = 0$ , means at  $x=0$ , 50% of the distribution lies to the left. (This is also the median)
    - >  $\text{Quantile}(0.95) = 1.65$



# Numerical Exploration

## > Relationships:

- `cov()`: covariances

$$\text{cov}(x, y) = E((x - \mu_x)(y - \mu_y))$$

- Interpretation: Expected value of the differences between x and y and their corresponding mean.
- E.g. if x is above it's mean when y is also above it's mean, then they will have a high covariance.
- Not bounded.



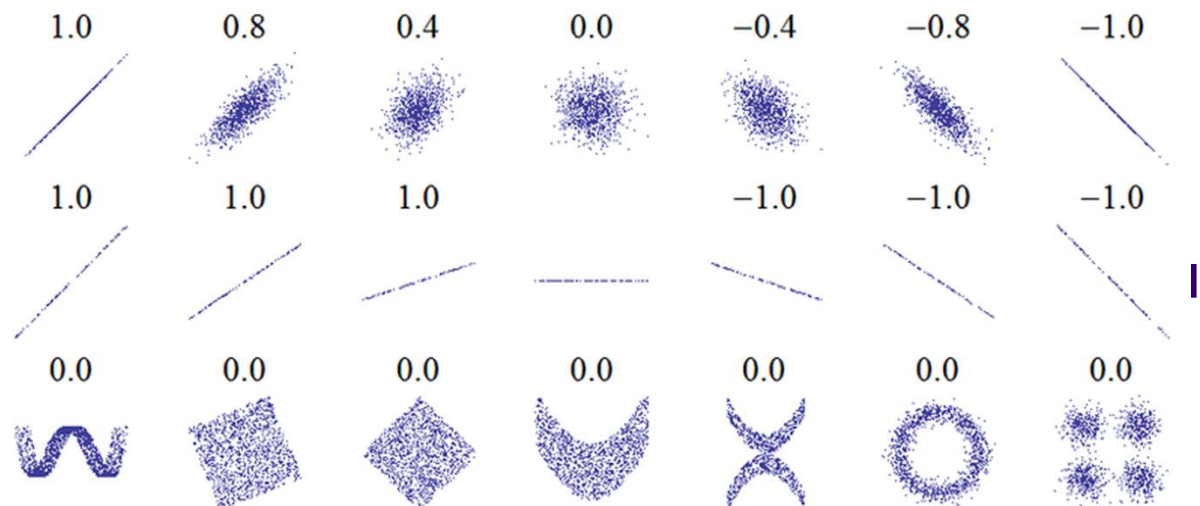
# Numerical Exploration

## > Relationships:

- `cor()`: correlations (pearsons)

$$\text{cor}(x, y) = \frac{E((x - \mu_x)(y - \mu_y))}{\sigma_x \sigma_y}$$

- Bounded between 0 and 1.
- **Does not indicate causation!!**



# Distribution Transformations

- > The purpose of transforming a variable is to make it easier to distinguish between values.
  - Most commonly we are looking to transform a distribution to be normal.
- > Common Transformations
  - Log-based:
    - >  $\text{Log}(x)$ ,  $\log(x+1)$ ,  $\log(x - \min(x) + 1)$
  - N-th Root based:
    - >  $X^{(1/n)}$
  - Any combination you can think of (remembering math rules).
- > We will cover normality tests in a later class.





# R Demo



Summary statistics





# Introduction to dplyr

## Data Wrangling Cheat Sheet:

<https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>

- > Data scientists spend most of their time on data munging or data wrangling
- > dplyr package provides a regular grammar for most data wrangling
- > Optimized for fast operations on data frames
- > Chain verbs (operators) for fast operations



Col1	Col2	Col3
2012	14	45
2013	13	76
2013	34	65
2014	23	47

```
library(dplyr)
df <- read.csv('data.csv',
header = TRUE,
stringsAsFactors = FALSE)
```

Col1	Col2	Col3
2012	12	45
2013	32	65
2013	34	65
2014	23	47

```
df <- filter(df, Col1 ==
2013)
```

Col1	Col3	Col3
2012	44	45
2013	76	76
2013	65	65
2014	47	47

```
df <- select(df, Col1, Col3)
```

Col1	Col2		Col3	Col4	
2012	14	14	45	45	59
2013	13	13	76	76	89
2013	34	34	65	65	99
2014	23	23	47	47	70

```
df <- mutate(df, Col4 = Col2 +  
Col3)
```

## Other useful dplyr verbs include:

```
df <- group_by(df, Col1)
```

```
df <- distinct(df, Col1)
```

```
df <- arrange(df, Col1)
```

```
df <- slice(df, 10:15)
```

```
df <- sample_frac(df, 0.5)
```

```
df <- sample_n(df, 500)
```

```
df <- summarize(df, m1 =  
mean(Col1))
```

Col1	Col2		Col3	Col4	
2012	13	14	76	45	89
2013	34	13	65	76	99
2013		34		65	
2014		23		47	

```
df <- df %>%
  filter(Col1 == 2013) %>%
  mutate(Col4 = Col2 + Col3)

iris %>% group_by(Species) %>%
  summarise(...)
```

# Simpsons Paradox

- > Slicing up data in different ways can create different results.
- > Generally arises in context of larger dataset with latent variable
- > <http://vudlab.com/simpsons/>
- > <http://www.math.grinnell.edu/~mooret/reports/SimpsonExamples.pdf>
- > **R Demo with dplyr**





# Exploratory data analysis


---

- > Explore the data with visualization
- > Understand the relationships in the data
- > Use multiple views of data
- > Aesthetics to project multiple dimensions
- > Conditioning to project multiple dimensions




# Seminal Book

John Tukey, Exploratory Data Analysis, 1977, Addison-Westley



# Seminal Book

**Visualizing Data**, William S. Cleveland, Hobart Press  
1993



**W**



# Views of data



- > Data contains complex relationships
- > Explore data with multiple views
- > Views show different aspects of the relationships
- > Different plots highlight different relationships



# Different plots for different views

- > Scatter
- > Scatter plot matrix
- > Line plots
- > Bar plots
- > Histograms
- > Box plots
- > Violin plots
- > Q-Q plots



# Visualization Aesthetics

Aesthetics expand dimensionality of projection

- > Color
- > Shape
- > Size
- > Transparency
- > Aesthetics specific to plot type
- > **Don't over do it!**



# Presenting charts

Charts must inform, not confuse!

- > Creating good visualization is iterative, and lots of work
- > Axis labels
- > Title
- > Legend
- > Large symbols and lines
- > Be sensitive to color blindness
- > No pie charts, please!
- > **Simplify, simplify, simplify!**

**W**

# Grammar of Graphics: ggplot2

## Grammar for building charts

### > Import library

```
library(ggplot2)
```

### > Define basic chart data and type

```
p1 = ggplot(df, aes(x = col1, y = col2,  
                    by = col3)) +  
      geom_plotttype(aes(asthetics))
```

### > Chain to add attributes

```
p1 + xlab('xlab') + ylab('ylab') +  
      ggtitle('The Title') +  
      other_attributes(...)
```

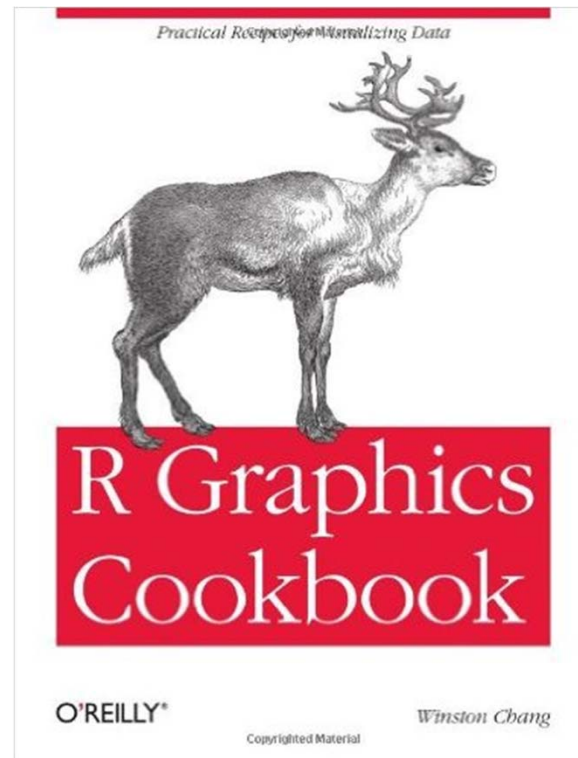




# ggplot2 resources

## ggplot2 cheat sheet

<https://www.rstudio.com/wp-content/uploads/2015/03/ggplot2-cheatsheet.pdf>



# W



# R Demo



Data Visualization



# Assignment

---

## Homework 1:

- Explore auto price data set.
- Write R program that shows/illustrates 3 key takeaways of your choosing from exploring the data.
- You should submit:
  - > **ONE R-script.**
  - > **One word document with 3 key points.** (example next page).



## Example Takeaway

- > The price of automobiles is dependent on feature  $x$ . The levels of feature  $x$  separate the following specific types of autos. Specifically these include..... The charts illustrate this relationship. Examining the chart shows.....



# Recommended Readings

---

- > An Introduction to Data Science, Chapters 3 and 9
- > Statistical Thinking for Programmers, Chapter 2

